

# Systems Security

## COMSM1500

# Change of plan

- Some students asked for lectures on the network-related security
  - Correspond to lab 4
- We will get back to concurrency-related vulnerabilities afterward
  - If any questions in the meantime
    - Come to office hours
    - Ask at the end of lectures

# Web Security

Client side

[bristol.ac.uk](http://bristol.ac.uk)



At the beginning of time...

# At the beginning of time...

- The web was simple

# At the beginning of time...

- The web was simple
- A server, a browser

# At the beginning of time...

- The web was simple
- A server, a browser
- The browser displayed the content sent by the server

# At the beginning of time

```
The World Wide Web project

WORLD WIDE WEB

The WorldWideWeb (W3) is a wide-area hypermedia[1] information retrieval
initiative aiming to give universal access to a large universe of documents.

Everything there is online about W3 is linked directly or indirectly to this
document, including an executive summary[2] of the project, Mailing lists[3] ,
Policy[4] , November's W3 news[5] , Frequently Asked Questions[6] .

    What's out there?[7]Pointers to the world's online information,
                        subjects[8] , W3 servers[9], etc.

    Help[10]           on the browser you are using

    Software           A list of W3 project components and their current
    Products[11]       state. (e.g. Line Mode[12] ,X11 Viola[13] ,
                        NeXTStep[14] , Servers[15] , Tools[16] , Mail
                        robot[17] , Library[18] )

    Technical[19]      Details of protocols, formats, program internals
                        etc

<ref.number>, Back, <RETURN> for more, or Help: █
```



# ... and now

```
The World Wide Web project

WORLD WIDE WEB

The WorldWideWeb (W3) is a wide-area hypermedia[1] information retrieval
initiative aiming to give universal access to a large universe of documents.

Everything there is online about W3 is linked directly or indirectly to this
document, including an executive summary[2] of the project, Mailing lists[3] ,
Policy[4] , November's W3 news[5] , Frequently Asked Questions[6] .

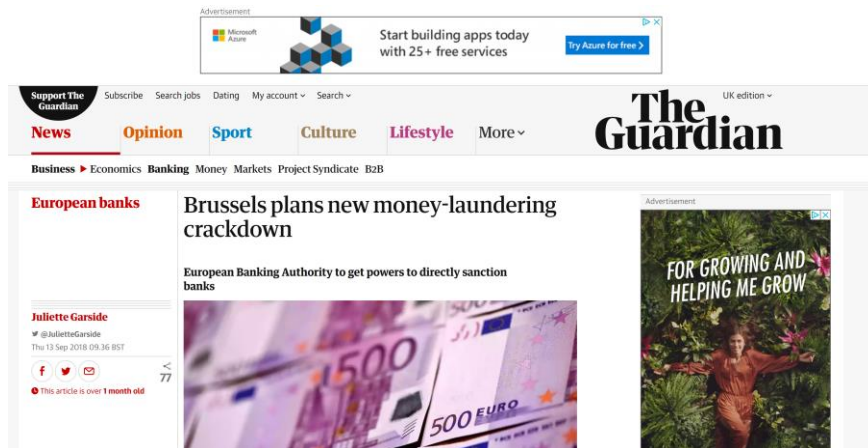
What's out there?[7]Pointers to the world's online information,
subjects[8] , W3 servers[9], etc.

Help[10]           on the browser you are using

Software           A list of W3 project components and their current
Products[11]       state. (e.g. Line Mode[12] ,X11 Viola[13] ,
                   NeXTStep[14] , Servers[15] , Tools[16] , Mail
                   robot[17] , Library[18] )

Technical[19]      Details of protocols, formats, program internals
                   etc

<ref.number>, Back, <RETURN> for more, or Help: █
```



bristol.ac.uk

# Feature rich browsers


- JavaScript
- DomModel
- AJAX
- Web sockets
- Multimedia
- Geolocation
- Many more features...

# Feature rich browsers


- JavaScript
- DomModel
- AJAX
- Web sockets
- Multimedia
- Geolocation
- Flash (we learned it was bad)
- Many more features...




# We learned



# We learned

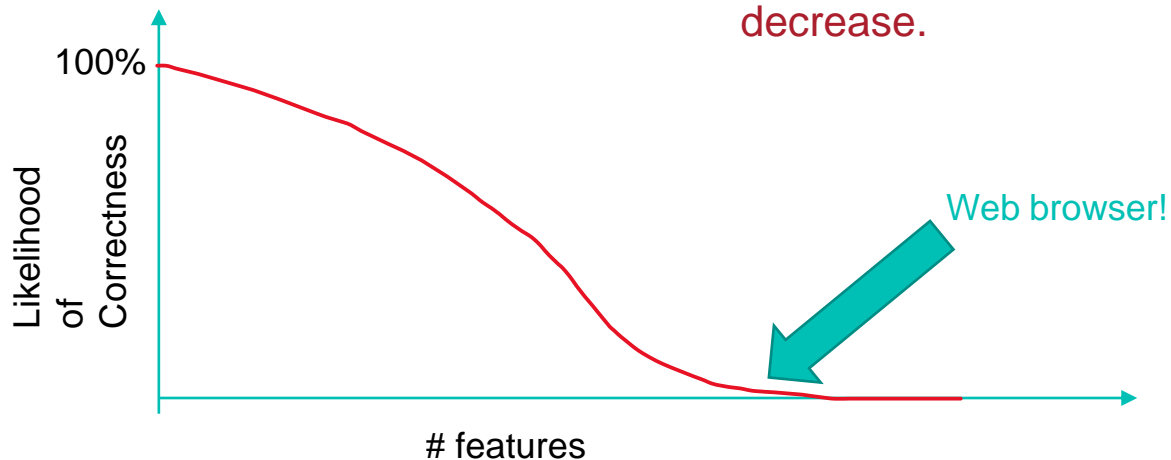


# We learned




# We learned

Homework/exam question:  
discuss why as the number  
of features increase, the  
correctness of a system  
decrease.



# What's on a website





# What's on a website

Advertisement from ads.com

The screenshot displays the homepage of The Guardian. At the top, there is a Microsoft Azure advertisement with the text "Start building apps today with 25+ free services" and a "Try Azure for free" button. Below this is the Guardian's navigation bar, which includes the "Support The Guardian" logo, links for "Subscribe", "Search jobs", "Dating", "My account", and "Search". The main navigation menu features categories: "News", "Opinion", "Sport", "Culture", "Lifestyle", and "More". The "UK edition" is also indicated. Below the navigation bar, a secondary menu lists "Business", "Economics", "Banking", "Money", "Markets", "Project Syndicate", and "BzB". The main content area features a headline "European banks" and "Brussels plans new money-laundering crackdown". Below the headline, it states "European Banking Authority to get powers to directly sanction banks". To the left of the article, there is a byline for "Juliette Garside" with her Twitter handle "@juliettegarside" and the date "Thu 13 Sep 2018 09:36 BST". Below the byline are social media sharing icons for Facebook, Twitter, and Email, along with a note "This article is over 1 month old" and a comment count of "77". The article's main image shows several 500 Euro banknotes. To the right of the article, there is another advertisement featuring a person in a field with the text "FOR GROWING AND HELPING ME GROW".

# What's on a website

Advertisement from ads.com

Analytics Library  
(e.g. from google.com)

bristol.ac.uk



# What's on a website

Advertisement from ads.com

Analytics Library  
(e.g. from google.com)

JQuery.js from  
Website.com



# What's on a website


Advertisement from ads.com

Analytics Library  
(e.g. from google.com)

JQuery.js from  
Website.com

HTML

bristol.ac.uk



# What's on a website

Advertisement from ads.com

Analytics Library  
(e.g. from google.com)

JQuery.js from  
Website.com

HTML  
+ inline JS

bristol.ac.uk



# What's on a website

Advertisement from ads.com

Analytics Library  
(e.g. from google.com)

JQuery.js from  
Website.com

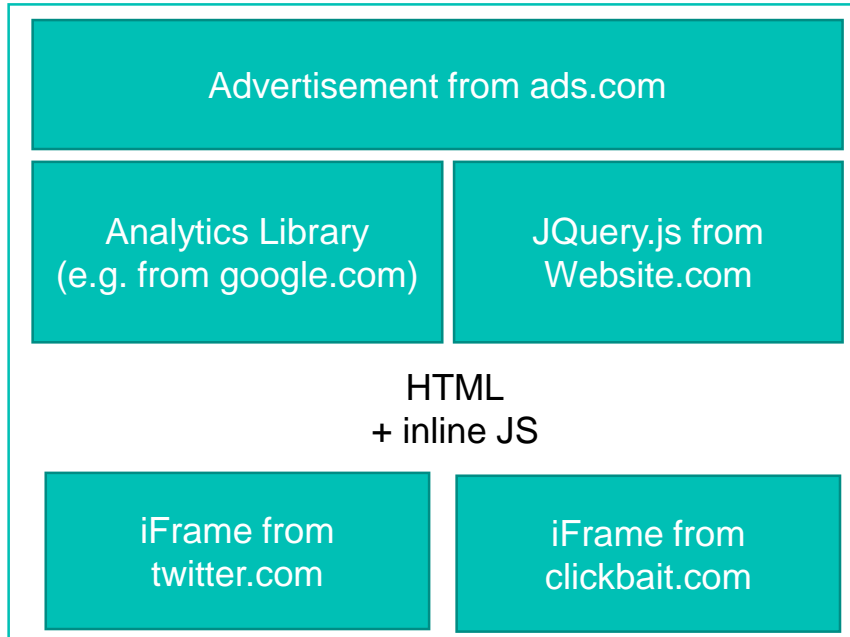
HTML  
+ inline JS

iFrame from  
twitter.com

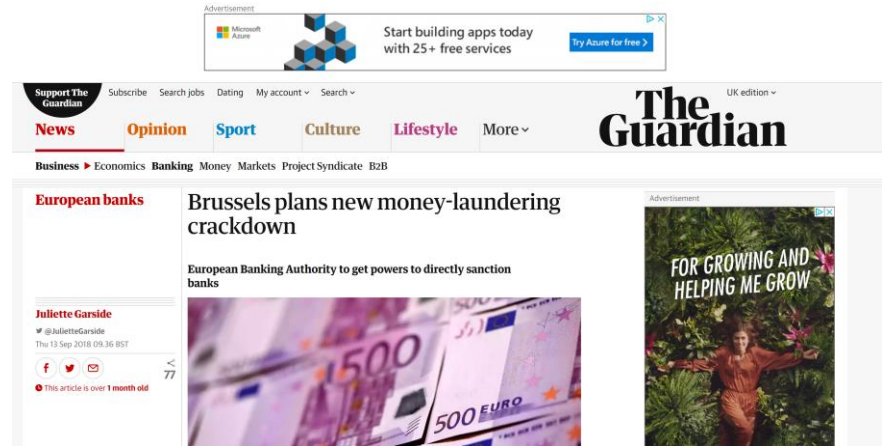
bristol.ac.uk



# What's on a website



bristol.ac.uk



# Same origin policy

How should we control how things interact?





# Same-origin policy

- Goal: two websites should not be able to tamper each others

# Same-origin policy

- Goal: two websites should not be able to tamper each others
  - A bit vague...
  - ... was easier in the early days
- Obviously bad: messing up with each other display

# Same-origin policy

- Goal: two websites should not be able to tamper each others
  - A bit vague...
  - ... was easier in the early days
- Obviously bad: messing up with each other display
- Obviously good: Google Map + Company Tracking

# Same-origin policy

- Goal: two websites should not be able to tamper each others
  - A bit vague...
  - ... was easier in the early days
- Obviously bad: messing up with each other display
- Obviously good: Google Map + Company Tracking
- More ambiguous: External JS library?

# Same-origin policy

- Strategy: each resource is assigned an origin. JS can only access resource from the same origin.

# Same-origin policy

- Strategy: each resource is assigned an origin. JS can only access resource from the same origin.
- Origin: scheme, domain, port
  - <http://foo.com/index.html>

# Same-origin policy

- Strategy: each resource is assigned an origin. JS can only access resource from the same origin.
- Origin: scheme, domain, port
  - <http://foo.com/index.html>
  - <https://foo.com/index.html>



Different scheme and port

# Same-origin policy

- Strategy: each resource is assigned an origin. JS can only access resource from the same origin.
- Origin: scheme, domain, port
  - <http://foo.com/index.html>
  - <https://foo.com/index.html>
  - [https://bar.com:8181/...](https://bar.com:8181/)



Different scheme and port






# Same-origin policy

Four fundamental ideas

# Same-origin policy

## Four fundamental ideas

- Each origin has client side resources
  - DOM tree (JS reflection of HTML page)
  - Cookies (to maintain states)
  - DOM storage (key/value store)
  - JS namespace
  - Display area



# Same-origin policy

Four fundamental ideas


- Each origin has client side resources
- Each frame get the origin of its URL

# Same-origin policy

Four fundamental ideas

- Each origin has client side resources
- Each frame get the origin of its URL
- Each JS scripts execute with the authority of its frame

# Same-origin policy



JS postMessage() interface  
(both side need to agree)

# Same-origin policy

## Four fundamental ideas

- Each origin has client side resources
- Each frame get the origin of its URL
- Each JS scripts execute with the authority of its frame
- Passive content get zero authority
  - Image
  - CCS

# Why rule 4?

- Mime sniffing attack

# Why rule 4?

- Mime sniffing attack
  - Old IE version trying to be helpful



# Why rule 4?

- Mime sniffing attack
  - Old IE version trying to be helpful
  - Idea: server may be misconfigured and misidentifies object type!

# Why rule 4?

- Mime sniffing attack
  - Old IE version trying to be helpful
  - Idea: server may be misconfigured and misidentifies object type!
  - A website link to an external image
  - IE would coerce it to for example JS script

# Why rule 4?

- Mime sniffing attack
  - Old IE version trying to be helpful
  - Idea: server may be misconfigured and misidentifies object type!
  - A website link to an external image
  - IE would coerce it to for example JS script
  - Get executed with the authority of the page it is in

# Why rule 4?

- Mime sniffing attack
  - Old IE version trying to be helpful
  - Idea: server may be misconfigured and misidentifies object type!
  - A website link to an external image
  - IE would coerce it to for example JS script
  - Get executed with the authority of the page it is in
- Browser are complex!
- Adding even well-meaning features can have unforeseen consequences!

# Why rule 4?

- Mime sniffing attack
  - Old IE version trying to be helpful
  - Idea: server may be misconfigured and misidentifies object type!
  - A website link to an external image
  - IE would coerce it to for example JS script
  - Get executed with the authority of the page it is in
- Browser are complex!
- Adding even well-meaning features can have unforeseen consequences!

# Frame/Window objects

- Frame get the origin of URL



OR

- A suffix of the original origin (set via document.domain)
  - e.g. X.Y.Z.com

# Frame/Window objects

- Get origin of the frame URL

OR

- A suffix of the original origin (set via document.domain)
  - e.g. X.Y.Z.com
  - Y.Z.com
  - Z.com
  - A.Y.Z.com 
  - .com 

# What could go wrong if not careful?

Hint UK





# What could go wrong if not careful?

- .co.uk
- .ac.uk
- etc...
- <https://publicsuffix.org/>

bristol.ac.uk

# Frame/Window objects

- Two frames can access each others if:
  - Both set document.domain to the same value
  - **Neither** has changed document.domain (and values match)
- Avoid attack from buggy subdomain

# Frame/Window objects

- Two frames can access each others if:
  - Both set document.domain to the same value
  - **Neither** has changed document.domain (and values match)
- Avoid attack from buggy subdomain
- Nice idea to get modules in different subdomains
  - e.g. login.foo.com; payment.foo.com etc...

# Cookies/AJAX/CCS

- Need to be subjected to similar origin rules
- Plugins are also a source of many nastiness
  - Used to be very trivial to write one to steal credit card number

# Some attacks



# Cross Site Scripting Attack



# Cross Site Scripting Attack

```
<html>  
<head>  
[...]  
</head>  
<body>  
[...]  
<p> [some user content] </p>  
</body>  
</html>
```

# Cross Site Scripting Attack


```
<html>  
<head>  
[...]  
</head>  
<body>  
[...]  
<p> [some user content] </p>  
</body>  
</html>
```






# Cross Site Scripting Attack

```
<html>  
<head>  
[...]  
</head>  
<body>  
[...]  
<p> [some user content] </p>  
</body>  
</html>
```



# Cross Site Scripting Attack

```
<html>  
<head>  
[...]  
</head>  
<body>  
[...]  
<p> </p><script>do.evill()</script></p>  
</body>  
</html>
```



# How to avoid this?



# Cross Site Scripting Attack


```
<html>
<head>
[...]
```

**DO SANITISE  
USER INPUT**

```
</head>
<body>
[...]
```

**<p> </p><script>do.evil()</script></p>**

```
</body>
</html>
```



# Cross Site Scripting Attack

Homework/exam question:  
discuss XSS attacks and  
how they can be easily prevented


```
<html>
<head>
[...]
```

**DO SANITISE  
USER INPUT**

```
</head>
<body>
[...]
```

`<p> </p><script>do.evil()</script></p>`

```
</body>
</html>
```



# Variation

- [https://mydomain.com/index.html?something=<script>do.evil\(\)</script>](https://mydomain.com/index.html?something=<script>do.evil()</script>)
  - Just get people to click on the link
- Could also get executed client side
  - *var url = new URL(url\_string);*
  - *var a = url.searchParams.get("something");*
- (little game at the end of the lecture)

# DNS exploit

[bristol.ac.uk](http://bristol.ac.uk)



# DNS exploit

- Rely heavily on DNS security
  - Based on domain name!



# DNS exploit

- Rely heavily on DNS security
  - Based on domain name!
- Approach:
  - Create a domain attacker.com

# DNS exploit

- Rely heavily on DNS security
  - Based on domain name!
- Approach:
  - Create a domain attacker.com
  - User visit attacker.com

# DNS exploit

- Rely heavily on DNS security
  - Based on domain name!
- Approach:
  - Create a domain attacker.com
  - User visit attacker.com
  - Browser generate DNS request to attacker.com

# DNS exploit

- Rely heavily on DNS security
  - Based on domain name!
- Approach:
  - Create a domain attacker.com
  - User visit attacker.com
  - Browser generate DNS request to attacker.com
  - Attacker response is very short lived (small TTL)

# DNS exploit

- Rely heavily on DNS security
  - Based on domain name!
- Approach:
  - Create a domain attacker.com
  - User visit attacker.com
  - Browser generate DNS request to attacker.com
  - Attacker response is very short lived (small TTL)
  - Attacker bind attacker.com to victim.com's IP

# DNS exploit

- Rely heavily on DNS security
  - Based on domain name!
- Approach:
  - Create a domain attacker.com
  - User visit attacker.com
  - Browser generate DNS request to attacker.com
  - Attacker response is very short lived (small TTL)
  - Attacker bind attacker.com to victim.com's IP
  - AJAX request intended for attacker.com actually goes to victim.com

# DNS exploit

- Rely heavily on DNS security
  - Based on domain name!
- Approach:
  - Create a domain attacker.com
  - User visit attacker.com
  - Browser generate DNS request to attacker.com
  - Attacker response is very short lived (small TTL)
  - Attacker bind attacker.com to victim.com's IP
  - AJAX request intended for attacker.com actually goes to victim.com
    - Attacker has code that executes within a company internal network
    - e.g. launch a port scan into a corporate network

# DNS exploit

- Rely heavily on DNS security
  - Based on domain name!
- Approach:
  - Create a domain attacker.com
  - User visit attacker.com
  - Browser generate DNS request to attacker.com
  - Attacker response is very short lived (small TTL)
  - Attacker bind attacker.com to some other IP
  - AJAX request intended for attacker.com actually goes to IP
    - You can start for example a port scan into a corporate network
- **FIX DNS RESOLUTION (TTL > some value)**



# DNS exploit

Homework/exam question:  
how an attacker could manipulate  
DNS TTL to perform port scan  
through javascript.

- Rely heavily on DNS security
  - Based on domain name!
- Approach:
  - Create a domain attacker.com
  - User visit attacker.com
  - Browser generate DNS request to attacker.com
  - Attacker response is very short lived (small TTL)
  - Attacker bind attacker.com to some other IP
  - AJAX request intended for attacker.com actually goes to IP
    - You can start for example a port scan into a corporate network
- **FIX DNS RESOLUTION (TTL > some value)**

# Exploiting page layout




# Exploiting page layout

- A pixel have no origin!


# Exploiting page layout

- A pixel have no origin!



# Exploiting page layout

- A pixel have no origin!



bristol.ac.uk


Attacker

Facebook



# Exploiting page layout

- A pixel have no origin!



# Not so private, private browsing



# Not so private, private browsing

- Client side privacy
  - Private browsing does not hide from network or server!



# Not so private, private browsing

- Client side privacy
  - Private browsing does not hide from network or server!
- Private browsing is very leaky, leave information in e.g.:
  - DNS cache
  - RAM artefacts in page swap, hibernation files

# Not so private, private browsing

- Client side privacy
  - Private browsing does not hide from network or server!
- Private browsing is very leaky, leave information in e.g.:
  - DNS cache
  - RAM artefacts in page swap, hibernation files
- Relatively trivial to recover such data with forensic tools!

# Not so private, private browsing

- Client side privacy
  - Private browsing does not hide from network or server!
- Private browsing is very leaky, leave information in e.g.:
  - DNS cache
  - RAM artefacts in page swap, hibernation files
- Relatively trivial to recover such data with forensic tools!
- Possible solution? Check the Veil paper (NDSS 2018)
  - on the course github ;)

# Not so private, private browsing

Homework/exam question:  
discuss why private browsing  
may create a false sense of  
privacy.

- Client side privacy
  - Private browsing does not hide from network or server!
- Private browsing is very leaky, leave information in e.g.:
  - DNS cache
  - RAM artefacts in page swap, hibernation files
- Relatively trivial to recover such data with forensic tools!
- Possible solution? Check the Veil paper (NDSS 2018)
  - on the course github ;)

# Conclusion

- Many features and they are used
- They are full of vulnerabilities
- However, we cannot walk back and provide something more secure
- We have to implement counter measure

- Google XSS game
  - <https://xss-game.appspot.com>

# Thank you

Office MVB 3.26

[bristol.ac.uk](http://bristol.ac.uk)

